

Graph Coloring Using State Space Search

Sandeep Dasgupta* Tanmay Gangwani† Mengjia Yan‡ Novi Singh §

December 10, 2014

1 Introduction

As a subject with many applications, graph coloring is one of the most studied NP-hard problems. Due to its time complexity graph coloring is an excellent candidate for implementation with a parallelized architecture, and as such has become our choice for our project. Of the many ways that graph coloring can be adapted for parallel programming there are two main approaches in literature: the iterative and state-space search methods. The iterative approach begins by dividing the vertices of the graph to be colored into different groups, each of which is assigned to a node in the parallelized system. In the context of Charm++ this would be akin to dividing the vertices among a set of chares. At this point each chare would independently color its assigned vertices, and once all vertices had been colored the chares would communicate with each other to see if each of their respective vertices had any neighbors on other chares that violate the constraints of the coloring problem. Once color conflicts have been assessed the chares would seek to mediate these problems through another round of independent coloring among the chares, a process that would be repeated until the graph has been successfully colored or shown to be uncolorable for a number of colors.

Unlike the iterative approach, state-space seeks to start with the initial graph and color it by creating child chares for each possible color assignment among the vertices. This process of coloration and child creation is repeated until a solution to the coloring problem is found or an exhaustive search has yielded no solution. By doing so the program explores all different possible colorations of the graph for a number of colors. State-space search has an advantage over the iterative approach in that it does not require the large number of messages sent by the methods chares to mediate coloring conflicts, nor does it have to deal with the complexity of mediation caused by highly connected subgraphs in the main graph. However, the state-space search does require the use of many more chares than the iterative approach and in doing so occurs higher memory overhead. In our opinion the message usage of the iterative approach outweighs the memory footprint of the state-space. Additionally, we chose to focus on finding the first solution given a user provided maximum of colors, k . Once found we keep on decrementing k until the graph become uncolorable. With this we can achieve the optimal coloring under the heuristics applied.

2 .ci file structure

Our implementation of state-space search is outlined in the structure of our .ci file. At the start of the program a root node is created by the main chare. This chare is responsible for graph

*Electronic address: sdasgup3@illinois.edu

†Electronic address: gangwan2@illinois.edu

‡Electronic address: mengjia@illinois.edu

§Electronic address: novi.singh94@gmail.com

preprocessing, namely precoloration, vertex removal, and independent subgraph detection (all will be discussed shortly). After preprocessing, the root selects a vertex to color, using a heuristic, and then spawns child nodes (chares) for each possible coloring of that vertex following a ordering based on another heuristic (called value ordering to be discussed later) and provides them with information regarding what has been colored already, and then waits for their response. The children of the root, and all subsequently created child nodes, also begin by preprocessing the graph, but only vertex removal and independent subgraph detection are applied. The response can be a success or failure and based on that the child states are merged with the parent to create an updated response which is propagated up to its parent. Each merge node has the intelligence to prune of redundant searches. For example, if a node which has spawned chares for independent graphs figures out that one of its subgraph is returning “failure of coloring”, then it ensures the the other subgraphs need not to be explored.

Grain size is the determining factor of the next step, causing the child to either spawn child nodes with value ordering if a determining factor is greater than the grain size, or to sequentially color the graph if it is equal to or less than it.

Chares are additionally grouped together, and result signals from the leaf nodes of the search are given higher privilege via the expedited tag. In doing so the problem of kill-chasing, that is the ineffective signaling of the end of the program due to competing the forces of kill messages and the spawning of new chares, is addressed.

3 Heuristics for State-Space Search

Due to the NP-hard nature of the problem the number of all possible states in the state-space search is exponential in size, and as such heuristics to reduce the number of states that are searched are necessary for any tractable computation to be possible. Some of the ideas are borrowed from [?].

- **Pre-Coloring**

The first of these heuristics is a preprocessing technique, namely the precoloring of the graph. By finding the 3-cliques in the graph we are able to preemptively assign them colors before the primary coloring algorithm begins. However, due to the extremely interconnected nature of dense graphs with large numbers of vertices, this technique is best applied to sparse graphs. As dense graphs tend to have many, many 3-cliques, this preprocessing technique often results in an almost entirely sequential coloring of the graph, may end up reporting un-colorability of the graph even though it is colorable. This is because the underlying algorithm to color is too naive, i.e. to put any possible color to a vertex such that it does not interfere with its neighbors.

- **Variable Ordering (Selecting a vertex to color next)**

An intuitive and simple heuristic guides our process for selecting the next vertex to be colored. By picking the vertex with the least available number of potential colors we can more effectively move throughout the graph providing a likely path towards the solution. This pairs well with our next pair of heuristics which are **impossibility testing** and **forced move** testing respectively. In the former, if after a vertex is colored if one of its neighbors has 0 available colors then the state for that coloring is not generated. The latter dictates that if the number of colors of the neighbor is brought down to 1, then the neighbor should be colored in that same step. This process is repeated recursively as well as coloring a neighbor might lead to another forced-move for its neighbor.

- **Vertex Removal**

The next applicable heuristic is vertex removal. By noting that a vertex with degree less than the number of available colors must always have a valid possible coloring, we can

remove such a vertex from the graph (pushing it onto a stack), color the rest of it, and then return the removed vertex to the graph and assign it a color. It is important to note that this removal can be performed recursive. That is, the removal of a vertex from the graph may lower the degree of a neighboring vertex to the point at which this heuristic can then be applied to it.

- **Detection of Independent Subgraphs**

This heuristic is to detect independent subgraphs, an $\mathcal{O}(V + E)$ operation, in the whole graph and then to run the coloring algorithm on each subgraph separately. By splitting the graph into two or more distinct pieces we often obtain large speedups as the problem is exponential in nature and evaluating multiple smaller subgraphs result in extremely fewer possible states as compared to evaluating a single, larger graph. This heuristic also leads to a changing of the structure of the state-space tree in that it transforms it into an AND-OR tree. In the single graph approach, either the left OR right (or at least one of many) has to find a solution to the coloring problem for a proper solution to be found, and failure in one subtree does not guarantee failure of the whole search. However, in the case of an AND-OR tree, when splitting the problem into multiple, independent subgraphs it becomes necessary for a solution to exist for all of the subgraphs (the left AND right subtrees have to yield solutions) as failure for one subgraph implies failure for the whole graph. Additionally, higher priority is given to those children of AND nodes which have less uncolored vertices with the hope that those will finish (get success or failure) faster. An important problem that our program has to deal with is the issue of kill-chasing. As previously stated we employ group chares for quickly stopping the spawning of new chares. Before creating any children a chore must obtain permission from the group chore, and if a solution is found then the group chore is notified immediately, halting the spawning of new state-space nodes. This notification is not only broadcasted to the group chore of the leaf that found the solution, but broadcasted to all group chares. This helps in greatly minimizing the traffic as opposed to the case when the information is sitting on a group (whose residing chore got the solution) and every chore is querying all the groups if any of them registered a success, before spawning its children.

Similarly, in the case of an AND node, if a failure in one of the subtrees of the node is found then a failure is immediately broadcasted to all group chares and spawning of new chares (for the other subtrees of that AND node) is efficiently halted with minimal messages.

- **Value Ordering and Prioritization**

Value ordering, a heuristic embedded in the structure of our .ci file, provides a lexicographical ordering of state-space search nodes to prioritize some children over others. This prioritization is based on the fact that the set of possible colors for a vertex need to be explored in the order of decreasing likelihood of getting the solution. The color that need to be explored first must be the one that affects its neighbor the least. To calculate this ordering we find for each coloring of the vertex, the sum of the number of remaining possible colors for all its neighbors. Then ranks are assigned to the colors in the decreasing order of the calculated sum, i.e. rank 0 is assigned to the color which corresponds to maximum sum. Based on this ranking bitvector are assigned and chares are spawned with this bitvector priorities which generates an ordering of chore spawning in the order of their rank, starting from 0.

Moreover the ordering is such that the left-most child of a node has a higher priority than the one to its right, resulting in a left-to-right, broom-stick sweep evaluation of the state-space tree. The major advantage that this provides is consistent and monotonic speedups which is otherwise not possible as the entire state space tree is expanded non-deterministically. Often this results in either the impossibility of a k-coloring, or a successful potential k-coloring, being detected deterministically between runs.

- **Grain size control**

Another important factor for limiting the number of chares that need to be created is grain size control. Our control hinges on the determining factor of the number of vertices that remain to be colored, and in the case of a number less than the grainsize the remainder of the graph is sequentially colored at the node. Picking a good grainsize is of paramount importance; too low and a huge number of chares are generated, incurring parallelization overhead, and a threshold that is too high results in few chares and an under-utilization of processing power. The sequential algorithm used in this control is a stack based, worklist approach that employs vertex removal and value ordering. Furthermore, this sequential coloring is **preemptable**, which is to say that after a set timeout value the continuation of the coloring is sent back into the same scheduling queue. By doing this we prevent the blocking of expedited messages, such as success or failure messages, and don't unnecessarily delay the termination of the program. An important note to the selection of grain size is that it had to be manually fine-tuned to find a good value as variations in graph structure and edge density would result in vastly different grain size optimums.

4 Experiments & Results

Table?? shows the generated random graphs along with their features which we used for our experimentations. Each of the graphs, G1 G6, were randomly generated by a python script. We varied the number of vertices, average edge densities, and maximum number of allowed colors, and additionally generated both truly random graphs and graph that were composed of one or more independent subgraphs. All of the graphs in the table were colorable save for two, and in the case of G5 and G6 we also conducted testing on detection of its explicitly generated independent subgraphs.

	Vertices	Edge Density	#Colors	Colorable
G1	300	12	9	Yes
G2	1000	5	5	Yes
G3	300	8	4	No
G4	500	7	6	Yes
G5	1000	6	6	Yes
G6	450	6	5	No

Table 1: Summary Of Graphs

The results of the execution of our program are summarized in the table??.

5 Seed Balancing

To further improve the execution time and processor utilization of our program we sought to employ a seed balancer. This balancer divides the creation of new chares across processors to even out the distribution of the workload. More specifically we tried both work stealing and "NeighborLB" as strategies for seed balancing. In the case of the former we found little difference in utilization compared to a non-load balanced execution, but in the latter we were able to note a major difference. In the Figure?? (b), generated using Projections, utilization remains higher for a longer period of time with NeighborLB enabled as compared to with any seed balancer, ?? (a). Furthermore, by looking at the distribution of the number of chares across each processor in figure??, we find that some processors have far fewer seed chares than their neighbors, something we attribute to the size of the workload of the chares on each processor being different.




Figure 1: Percentage Utilization on graph G3 with grainsize 260, timeout = 10s (a) Without Seed balancer (b) With "Neighbor" Seed Balancer




Figure 2: Distribution of seeds on each on 12 PEs. The Seed Balancer is able to well distribute the seeds based on the PE loads. It is intuitive that PE 7 must be highly overloaded so that it has been given the least number of seeds.

	Vertices	Grain Size	Timeout	Priority Bits	Value Ordering	Sub-graph Detection	Execution Time (s)	#Chares
G1	300	290	10	Enable	Enable	Enable	18	1220
				Disable			31	189105
G3	300	260	10	Enable			27	186329
		280					16	124601
		290					7	2681
		300					82	1
G2	1000	960	10	Enable	Enable	Enable	41	235
			30				214	585
G4	500	480	10	Enable	Enable	Enable	102	642
					Disable		133	669
G5	1000	960	5	Enable	Enable	Enable	0.01	4
						Disable	35	2610
G6	450	410	5	Enable	Enable	Enable	0.023997	4
						Disable	5.769123	49646

Table 2: Results with different features enabled at a time. All runs are made on 12 PEs on taub.campuscluster.illinois.edu, except the last one which is done on 16 PEs.




Figure 3: Graph used is un-colorable G3, with 4 colors (vertices 300, Edge Density 8).

6 Core Scaling

For an un-colored graph all states need to be explored for a failure to be found and more PEs allows for a great parallelization of the search. As shown in Figure ??, the core scaling of our program was shown to be consistent for such graphs in the sense that more processors results in a smaller execution time. However, diminishing results are noted as the number of PEs increases, a fact that is attributed to the marginal reduction of time to reach a leaf node given our usage of node prioritization. Past a certain point the amount of work per PE to reach a solution containing leaf node does not decrease much as the work becomes adequately distributed rather quickly. The reason for

7 Conclusion

Our project sought to employ as many heuristics as possible to mitigate the exponential nature of the graph coloring problem, and this desire combined with experimentation of other execution

parameters has allowed us to create an effective program for graph coloring. Many of our difficulties lay in fine-tuning our heuristics to work with the large variation in graph structure that could occur, as well as working with the non-deterministic effects in the exploration of the state-space tree. In the future this work could be expanded upon to more clearly express the impact of each heuristic on a graph, and to change the applications of each heuristic automatically depending on the characteristics of the graph. That being said, we believe we have provided a solid groundwork for such work to be started.

8 References